

Chapters 6 and 7

Introduction to Pandas and Access Operations



- Pandas is an open-source Python library for data analysis that was originally developed by Wes McKinney in 2008. The name Pandas stands for “Python Data Analysis Library” and is derived from the term “panel data,” which is an econometrics term for multidimensional structured data sets.
- Pandas gives Python the ability to work with spreadsheet-like data and is built on top of NumPy.
- Pandas introduces two new data types to Python: `Series` and `DataFrame`.

- The DataFrame is a two-dimensional data structure where data is aligned in a tabular fashion in rows and columns. Essentially, a Pandas DataFrame is an in-memory representation of an Excel spreadsheet in Python. A Pandas Series is a single column of a DataFrame.
- Each column (Series) of a DataFrame has to be the same type (just like a NumPy array), whereas each row can contain mixed types.
- A simple DataFrame is illustrated below ([source](#)):

	Name	Team	Number	Position	Age
0	Avery Bradley	Boston Celtics	0.0	PG	25.0
1	John Holland	Boston Celtics	30.0	SG	27.0
2	Jonas Jerebko	Boston Celtics	8.0	PF	29.0
3	Jordan Mickey	Boston Celtics	NaN	PF	21.0
4	Terry Rozier	Boston Celtics	12.0	PG	22.0
5	Jared Sullinger	Boston Celtics	7.0	C	NaN
6	Evan Turner	Boston Celtics	11.0	SG	27.0

- At a basic level, Pandas DataFrame objects can be thought of as enhanced versions of two-dimensional NumPy arrays in which the rows and columns are identified with labels rather than simple integer indices.

Installing and Importing Pandas

- The Pandas package is not part of the standard Python release and may need to be installed separately. If you use pip, you can install Pandas with:

```
pip3 install pandas
```

- Once Pandas is installed, you need to import the Pandas library as follows:

```
import pandas as pd
```

- It is common practice to import Pandas with the alias 'pd'.

The Pandas Series Object

- The Series object is a one-dimensional array of indexed data. It can be created from a list or array as follows.

```
import pandas as pd
data = pd.Series([10, 30, -6, 9])
print(data)
```

```

0    10
1    30
2    -6
3     9

```

- As you can see, the `Series` object contains an array of data (of any NumPy data type) with associated indices (which can be numbers, strings, or any other data types). We can access these with the `values` and `index` attributes.

```

print(data.values)
print(data.index)

```

```

[10 30 -6  9]
RangeIndex(start=0, stop=4, step=1)

```

- It appears that a `Series` is interchangeable with a one-dimensional NumPy array. However, it comes with some additional functionality and can be thought of as a column of a spreadsheet. The `Series` can have a name, and most importantly, it has an explicitly defined index associated with the values.
- By default, the `Series` will be indexed from 0 to `n` where `n = size-1`.

```

age = pd.Series([10,13,9,16],index=['Bill','Peixin','Carlo','Thahn'],name='kids_ages')
print(age)

```

```

Bill      10
Peixin    13
Carlo      9
Thahn     16
Name: ages, dtype: int64

```

- We can access the elements of a `Series` by either the index *position* or *label*:

```

>>> age[1]
13

>>> age['Peixin']
13

```

- You can also construct a `Series` from a dictionary as illustrated in the example below ([source](#)):

```

population_dict = {'California': 38332521,
                   'Texas': 26448193,
                   'New York': 19651127,
                   'Florida': 19552860,
                   'Illinois': 12882135}
population = pd.Series(population_dict)
print(population)

```

```
California    38332521
Texas        26448193
New York     19651127
Florida      19552860
Illinois     12882135
dtype: int64
```

- Here are some of the common attributes of a Series object:

Attribute	Returns
name	The name of the Series object
dtype	The data type of the Series object
shape	Dimensions of the Series object in a tuple of the form (# of rows,)
index	The Index object that is part of the Series object
columns	The name of the columns (as an Index object)
values	The data in the Series object

- Note that the Index class makes the Series class more powerful than a NumPy array because it gives us row labels.
- Many of the methods and functions that operate on a NumPy array will also operate on a Series. For example:

```
>>> age.mean()
12.0
>>> age.max()
16
>>> age.std()
3.1622776601683795
```

- The Pandas.Series [documentation](#) contains more information on how to create a Series object and the full list of attributes and methods that are available.

The Pandas DataFrame Object

- The DataFrame object is a two-dimensional array with both flexible row indices and flexible column names. This is illustrated graphically below ([source](#)).

The diagram illustrates a Pandas DataFrame with 7 rows and 9 columns. The columns are labeled: Name, Team, Number, Position, Age, Height, Weight, College, and Salary. The rows are indexed from 0 to 6. The data is as follows:

	Name	Team	Number	Position	Age	Height	Weight	College	Salary
0	Avery Bradley	Boston Celtics	0.0	PG	25.0	6-2	180.0	Texas	7730337.0
1	John Holland	Boston Celtics	30.0	SG	27.0	6-5	205.0	Boston University	NaN
2	Jonas Jerebko	Boston Celtics	8.0	PF	29.0	6-10	231.0	NaN	5000000.0
3	Jordan Mickey	Boston Celtics	NaN	PF	21.0	6-8	235.0	LSU	1170960.0
4	Terry Rozier	Boston Celtics	12.0	PG	22.0	6-2	190.0	Louisville	1824360.0
5	Jared Sullinger	Boston Celtics	7.0	C	NaN	6-9	260.0	Ohio State	2569260.0
6	Evan Turner	Boston Celtics	11.0	SG	27.0	6-7	220.0	Ohio State	3425510.0

Annotations in the diagram include: 'Column names' pointing to the header row, 'Columns axis=1' pointing to the column headers, 'Index label' pointing to the row indices, 'Index axis=0' pointing to the row indices, 'Missing value' pointing to the NaN values in the 'Number' and 'Age' columns, and 'Data' pointing to the numerical data cells.

- Most of the time, a Pandas DataFrame will be created by importing data from existing storage, such as a CSV file, an SQL database, or an Excel spreadsheet.

Creating a DataFrame

- A Pandas DataFrame can be constructed in a number of ways, but we will demonstrate constructing one from a dictionary of lists (DoL).

```
dataDoL={'name':['Bella','Charlie','Lucy','Coooper','Max','Stella','Bernie'],
        'breed':['Labrador','Poodle','ChowChow','Schnauzer','Labrador',
                 'Chihuahua','St.Bernard'],
        'color': ['Brown','Black','Brown','Gray','Black','Tan','White'],
        'height_cm':[56,43,46,49,59,18,77],
        'weight_kg':[24,24,24,17,29,2,74],
        'date_of_birth':['2013-07-01','2016-09-16','2014-08-25',
                        '2011-12-11','2017-01-28','2015-04-20',
                        '2018-02-27']}

dogs = pd.DataFrame(dataDoL)
print(dogs)
```

	name	breed	color	height_cm	weight_kg	date_of_birth
0	Bella	Labrador	Brown	56	24	2013-07-01
1	Charlie	Poodle	Black	43	24	2016-09-16
2	Lucy	ChowChow	Brown	46	24	2014-08-25
3	Coooper	Schnauzer	Gray	49	17	2011-12-11
4	Max	Labrador	Black	59	29	2017-01-28
5	Stella	Chihuahua	Tan	18	2	2015-04-20
6	Bernie	St.Bernard	White	77	74	2018-02-27

- Pandas does allow for row labels, but since we have not specified the labels, Pandas has provided a set of integers (0–6) as the row labels.
- Now the columns are in the order we originally specified.
- The follow are some commonly used DataFrame attributes.

Attribute	Returns
dtypes	The data type of each column
shape	Dimensions of the DataFrame object in a type of the form (# of rows, # of columns)
index	The Index object along the rows of the DataFrame.
columns	The name of the columns (as an Index object)
values	The data in the DataFrame object
empty	Check if the DataFrame object is empty.

- We can specify row labels by assigning values to the index attribute.

- As an example, consider the following data set of country-based indicators. We will assign the country codes as the row labels.

```
indicatorDoL = {
    'country': ['Canada', 'China', 'India',
                'Russia', 'United States', 'Vietnam'],
    'pop': [36.26, 1378.66, 1324.17, 144.34, 323.13, 94.59],
    'gdp': [1535.77, 11199.15, 2263.79, 1283.16, 18624.47,
            205.28],
    'life': [82.30, 76.25, 68.56, 71.59, 78.69, 76.25],
    'cell': [30.75, 1364.93, 1127.81, 229.13, 395.88, 120.60]}

codes = pd.Index(['CAN', 'CHN', 'IND', 'RUS', 'USA', 'VNM'],
                  name='code')

indicators = pd.DataFrame(indicatorDoL, index=codes)

print(indicators)
```

	country	pop	gdp	life	cell
code					
CAN	Canada	36.26	1535.77	82.30	30.75
CHN	China	1378.66	11199.15	76.25	1364.93
IND	India	1324.17	2263.79	68.56	1127.81
RUS	Russia	144.34	1283.16	71.59	229.13
USA	United States	323.13	18624.47	78.69	395.88
VNM	Vietnam	94.59	205.28	76.25	120.60

- Below we demonstrate accessing the values, columns, and index of a DataFrame.

```
print(dogs.values)
print(dogs.columns)
print(dogs.index)
```

```
[['Bella' 'Labrador' 'Brown' 56 24 '2013-07-01']
 ['Charlie' 'Poodle' 'Black' 43 24 '2016-09-16']
 ['Lucy' 'ChowChow' 'Brown' 46 24 '2014-08-25']
 ['Cooper' 'Schnauzer' 'Gray' 49 17 '2011-12-11']
 ['Max' 'Labrador' 'Black' 59 29 '2017-01-28']
 ['Stella' 'Chihuahua' 'Tan' 18 2 '2015-04-20']
 ['Bernie' 'St.Bernard' 'White' 77 74 '2018-02-27']]

Index(['name', 'breed', 'color', 'height_cm', 'weight_kg', 'date_of_birth'],
      dtype='object')

RangeIndex(start=0, stop=7, step=1)
```

Importing a CSV File and Exploring a DataFrame

- Recall that a **CSV (comma separated values)** files are one of the most popular file formats for importing and exporting tabular data such as spreadsheets and databases. For each row, the column information is separated with a comma. However, the comma is not the only delimiter. Some files are delimited by a tab (TSV file) or even a semicolon. The main

reason why CSVs are a preferred data format when collaborating and sharing data is because any program can open this kind of data structure, including a text editor.

- To import a CSV file, we can use the Pandas `read_csv(.)` function. We will explore the full functionality of this function later, but we will use it in its most basic form here.
- Download the `gapminder.csv` file from Teams and put it in your working directory. This is a data set prepared by Jennifer Bryan and is a subset of the [gapminder](#) teaching package for R. We can then import the file as follows.

```
df = pd.read_csv("gapminder.csv")
print("The data type of df is " + str(type(df)))

print('The shape of the dataframe is ' + str(df.shape))
print(df.head())
```

```
The data type of df is <class 'pandas.core.frame.DataFrame'>
```

```
The shape of the dataframe is (1704, 6)
```

	country	continent	year	lifeExp	pop	gdpPercap
0	Afghanistan	Asia	1952	28.801	8425333	779.445314
1	Afghanistan	Asia	1957	30.332	9240934	820.853030
2	Afghanistan	Asia	1962	31.997	10267083	853.100710
3	Afghanistan	Asia	1967	34.020	11537966	836.197138
4	Afghanistan	Asia	1972	36.088	13079460	739.981106

- We used the Python `type` function to verify that `df` is a Pandas `DataFrame` and we used the `shape` *attribute* of the `DataFrame` object to obtain the number of rows (1704) and the number of columns (6).
- Finally, to get a sense of the `DataFrame` contents we used the Pandas `head(.)` function that returns the first few rows of the `DataFrame`.
- We can use the `info(.)` method to display the names of the columns, the data types they contain, and whether they have any missing values.

```
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1704 entries, 0 to 1703
Data columns (total 6 columns):
country      1704 non-null object
continent    1704 non-null object
year         1704 non-null int64
lifeExp      1704 non-null float64
pop          1704 non-null int64
gdpPercap    1704 non-null float64
dtypes: float64(2), int64(2), object(2)
```

- Note that each column has a single data type, but they do not all share the same type.

- The `describe()` method computes some summary statistics for numerical columns, including the mean, standard deviation, and the five-number summary. Note that `count` is the number of non-missing values in each column.

```
print(df.describe())
```

	year	lifeExp	pop	gdpPercap
count	1704.000000	1704.000000	1.704000e+03	1704.000000
mean	1979.500000	59.474439	2.960121e+07	7215.327081
std	17.265333	12.917107	1.061579e+08	9857.454543
min	1952.000000	23.599000	6.001100e+04	241.165877
25%	1965.750000	48.198000	2.793664e+06	1202.060309
50%	1979.500000	60.712500	7.023596e+06	3531.846989
75%	1993.250000	70.845500	1.958522e+07	9325.462346
max	2007.000000	82.603000	1.318683e+09	113523.132900

Access Operations

- Access operations are those that *read* or *query* data values out of a table. The most common access operations are:
 1. **Single-element access:** We may want to access a single data value (at the intersection of a row and column).
 2. **Column access:** We seek a subset of a data frame that is based on one or more columns, which we call a *projection* of the desired columns. For example, we may want a subtable obtained by projecting the `year` and `lifeExp` columns (with all rows included).
 3. **Single row:** We may wish to select a single row, containing the values of all columns in that row.
 4. **Multiple rows:** We want a subset of the data consisting of all the columns, but a limited set of the rows.
 5. **Subset of rows and columns:** The most general form of access operation would allow us to both project a subset of the columns *and* filter for a particular subset of the rows.
- We can **select a single column** of a `DataFrame` by using the name of the `DataFrame`, followed by square brackets with a column name inside. When we project a single column in this way, we obtain a `Series` as the projected column.

```
Name = dogs['name']
print(Name)
print(type(Name))
```

0	Bella
1	Charlie
2	Lucy
3	Coooper
4	Max
5	Stella


```
6      Bernie
Name: name, dtype: object

<class 'pandas.core.series.Series'>
```

- We can also project a single column by making the *column name* an *attribute* of the DataFrame object as demonstrated below.

```
print(dogs.name)
```

```
0      Bella
1     Charlie
2       Lucy
3     Cooper
4       Max
5     Stella
6     Bernie
Name: name, dtype: object
```

- To **select (project) multiple columns**, you need two pairs of square brackets. The outer brackets are responsible for subsetting the DataFrame, and the inner square brackets are creating a list of column names to subset.

```
subtable = dogs[['height_cm', 'breed']]
print(subtable)
type(subtable)
```

```
   height_cm  breed
0         56  Labrador
1         43   Poodle
2         46  ChowChow
3         49  Schnauzer
4         59  Labrador
5         18  Chihuahua
6         77 St.Bernard

pandas.core.frame.DataFrame
```

- First, note that we requested the columns in an order different from their order in the DataFrame and this requested order was respected in the result. Second, note that the type returned is a DataFrame.
- There are a number of ways to **access rows**. One of the easiest methods is to use slicing notation. Keep in mind that row (and column) numbers start at 0 and up to the number of rows (or columns) minus one. The syntax for a slice is:

```
start : end [:stride]
```

- Recall our indicators DataFrame:

	country	pop	gdp	life	cell
code					
CAN	Canada	36.26	1535.77	82.30	30.75
CHN	China	1378.66	11199.15	76.25	1364.93
IND	India	1324.17	2263.79	68.56	1127.81

RUS	Russia	144.34	1283.16	71.59	229.13
USA	United States	323.13	18624.47	78.69	395.88
VNM	Vietnam	94.59	205.28	76.25	120.60

```
subrows = indicators[3:5]
print(subrows)
```

	country	pop	gdp	life	cell
code					
RUS	Russia	144.34	1283.16	71.59	229.13
USA	United States	323.13	18624.47	78.69	395.88

- The output above corresponds to the rows indexed by 3 and 4 (since 5 is not included). We can select rows starting at the beginning and proceeding up to, but not including index 2, as follows.

```
print(indicators[:2])
```

	country	pop	gdp	life	cell
code					
CAN	Canada	36.26	1535.77	82.30	30.75
CHN	China	1378.66	11199.15	76.25	1364.93

- We can also access rows using the labels (index) attribute within the slice.

```
print(indicators['IND':'USA'])
```

	country	pop	gdp	life	cell
code					
IND	India	1324.17	2263.79	68.56	1127.81
RUS	Russia	144.34	1283.16	71.59	229.13
USA	United States	323.13	18624.47	78.69	395.88

A stride of -1
reverses the order

```
print(indicators['VNM':'RUS':-1])
```

	country	pop	gdp	life	cell
code					
VNM	Vietnam	94.59	205.28	76.25	120.60
USA	United States	323.13	18624.47	78.69	395.88
RUS	Russia	144.34	1283.16	71.59	229.13

Row Selection by Condition

- Another common method to select rows is to use a logical condition inside of square brackets to filter against. For example, we can find all the dogs whose height is greater than 50 centimeters as follows.

```
print(dogs[dogs['height_cm'] > 50])
```

	name	breed	color	height_cm	weight_kg	date_of_birth
0	Bella	Labrador	Brown	56	24	2013-07-01
4	Max	Labrador	Black	59	29	2017-01-28
6	Bernie	St.Bernard	White	77	74	2018-02-27

- Let's tease out this statement and look at what `dogs['height_cm'] > 50` returns.

```
print(dogs['height_cm'] > 50)
```

```
0      True
1     False
2     False
3     False
4      True
5     False
6      True
Name: height_cm, dtype: bool
```

- This statement returns a `Series` with a `dtype` of `bool`. Therefore, we can subset values using labels and indices, but also by supplying a vector of Boolean values.
- We can also subset rows based on text data. In the example below, we use the double equal sign in the logical condition to filter the dogs that are Labradors.

```
print(dogs[dogs['breed']=='Labrador'])
```

	name	breed	color	height_cm	weight_kg	date_of_birth
0	Bella	Labrador	Brown	56	24	2013-07-01
4	Max	Labrador	Black	59	29	2017-01-28

- We can also subset based on dates. Below we filter all the dogs born before 2015. Notice that the dates are in quotes and are written in the *international standard date* format.

```
print(dogs[dogs["date_of_birth"] > "2015-01-01"])
```

	name	breed	color	height_cm	weight_kg	date_of_birth
1	Charlie	Poodle	Black	43	24	2016-09-16
4	Max	Labrador	Black	59	29	2017-01-28
5	Stella	Chihuahua	Tan	18	2	2015-04-20
6	Bernie	St.Bernard	White	77	74	2018-02-27

- More tips for working with dates: Pandas supports datetime format for dates to convert a string date column into datetime format and do manipulations or calculations on it. For example, to find the age of each dog in years, you can run

```
import datetime as dt # need to install it if it is not installed
import numpy as np # for rounding
dogs['date_of_birth'] = pd.to_datetime(dogs['date_of_birth'])
current_date = dt.date.today()
dogs['current_date'] = pd.to_datetime(current_date)
np.round((dogs['current_date'] - dogs['date_of_birth']).dt.days/365,2)
```

- To **subset rows that meet multiple conditions**, you can combine conditions using logical operations, such as the “&” operator as seen below.

```
is_lab = dogs['breed'] == 'Labrador'
is_brown = dogs['color'] == 'Brown'
print(dogs[is_lab & is_brown])
```

	name	breed	color	height_cm	weight_kg	date_of_birth
0	Bella	Labrador	Brown	56	24	2013-07-01

- Note that we could have also done this in one line of code, but you would need to add parentheses around each condition.
- If you want to filter on multiple values of a categorical variable, the easiest way is to use the `isin(.)` method. This takes in a list of values to filter for. Below, we check if the color of a dog is black or brown.

```
is_black_or_brown = dogs[dogs['color'].isin(['Black', 'Brown'])]
print(dogs[is_black_or_brown])
```

	name	breed	color	height_cm	weight_kg	date_of_birth
0	Bella	Labrador	Brown	56	24	2013-07-01
1	Charlie	Poodle	Black	43	24	2016-09-16
2	Lucy	ChowChow	Brown	46	24	2014-08-25
4	Max	Labrador	Black	59	29	2017-01-28

Selection using `loc` and `iloc`

- For any Pandas DataFrame, we can use `.loc[]` and `.iloc[]` to perform practically any data selection. **Note that `loc` is label-based, while `iloc` is integer index based.** The syntax is as follows.

```
dataframe.loc[rowspec, colspec]
```

```
dataframe.iloc[rowspec, colspec]
```

Inside the access operator, *rowspec* and *colspec* are row and column *specifiers*.

- We provide some examples of using `loc` and `iloc` below.

```
print(indicators.loc['VNM', 'gdp'])
```

```
| 205.28
```

Here we access a single element, and the result is a scalar with the data type of the entry.

```
print(indicators.loc[:, 'pop'])
```

```
code
CAN      36.26
CHN    1378.66
IND    1324.17
RUS     144.34
USA     323.13
VNM       94.59
```

Here we are accessing a single column, `pop`, and all the associated rows (`:`).

```
| Name: pop, dtype: float64
```

```
print(indicators.iloc[1:3, :])
```

	country	pop	gdp	life	cell
code					
CHN	China	1378.66	11199.15	76.25	1364.93
IND	India	1324.17	2263.79	68.56	1127.81

Here we use `iloc` to access the rows indexed from 1 to 3 (but not including 3) and all the associated columns.

```
print(indicators.loc['USA', ['country', 'life', 'cell']])
```

country	United States
life	78.69
cell	395.88

Name: USA, dtype: object

Here we access a subset of the columns (`country`, `life`, and `cell`) for a single row (`USA`). Note that this returned a `Series` (and thus is now a column!).

```
print(indicators.iloc[4, [0, 3, 4]])
```

country	United States
life	78.69
cell	395.88

Name: USA, dtype: object

Here we access the same information as the last example, but we do so using the row and column indices (with `iloc`) as opposed to the labels.

```
print(indicators.loc[ indicators.gdp < 2000, ['country', 'pop']])
```

	country	pop
code		
CAN	Canada	36.26
RUS	Russia	144.34
VNM	Vietnam	94.59

Here we access the two columns `country` and `pop` for the rows for which the `gdp` is less 2000.

Chapters 8

Advanced Operations in Pandas

Aggregating a Single Pandas Series

- Data analysis typically requires us to summarize data from a subset of a variable, called *aggregating*.
- The simplest form of aggregation occurs when we have a single Series (often a single column projection of a DataFrame). The most common aggregation functions/methods are provided in the table below.

Table 8.1 Aggregation function/methods

Method	Description
<code>mean()</code>	Arithmetic mean, not including missing values
<code>median()</code>	Value occurring halfway through the population, omitting missing values
<code>sum()</code>	Arithmetic sum of the non-missing values
<code>min()</code>	Smallest value in the set
<code>max()</code>	Largest value in the set
<code>nunique()</code>	Number of unique values in the set
<code>size()</code>	Size of the set
<code>count()</code>	Number of non-missing values in the set

- Two other handy methods are in the table below.

Table 8.2 Aggregation to Series methods

Method	Description
<code>unique()</code>	Construct a subset Series consisting of the unique values from the source Series
<code>value_count()</code>	Construct a Series of integers capturing the number of times each unique value is found in a source Series; the Index of the new Series consists of the unique values from the source Series

- We demonstrate some of these methods below on the dogs DataFrame from the previous chapter:

	name	breed	color	height_cm	weight_kg	date_of_birth
0	Bella	Labrador	Brown	56	24	2013-07-01
1	Charlie	Poodle	Black	43	24	2016-09-16
2	Lucy	ChowChow	Brown	46	24	2014-08-25
3	Coooper	Schnauzer	Gray	49	17	2011-12-11
4	Max	Labrador	Black	59	29	2017-01-28
5	Stella	Chihuahua	Tan	18	2	2015-04-20
6	Bernie	St.Bernard	White	77	74	2018-02-27

```

avg_weight = dogs['weight_kg'].mean()
max_height = dogs['height_cm'].max()
min_weight = dogs['weight_kg'].min()
numColors = dogs['color'].nunique()

print('Mean weight = ', avg_weight)
print('Max height = ', max_height)
print('Min weight = ', min_weight)
print('Number of different colors = ', numColors)

```

```

Mean weight = 27.714285714285715
Max height = 77
Min weight = 2
Number of different colors = 5

```

Aggregating a DataFrame

- We are now interested in grouping rows into discrete partitions, aggregating each partition into a row of computed values, and combining the result.
- In Pandas, we use the `agg()` method of a `DataFrame` to perform this general form of the aggregation operation. The argument to `agg()` must convey, for each column of the original `DataFrame`, what aggregation function or functions to perform.
- Recall our indicators `DataFrame`:

	country	pop	gdp	life	cell
code					
CAN	Canada	36.26	1535.77	82.30	30.75
CHN	China	1378.66	11199.15	76.25	1364.93
IND	India	1324.17	2263.79	68.56	1127.81
RUS	Russia	144.34	1283.16	71.59	229.13
USA	United States	323.13	18624.47	78.69	395.88
VNM	Vietnam	94.59	205.28	76.25	120.60

- We demonstrate on the indicators `DataFrame` below, performing two aggregations on the `life` column (min and median), and just the median aggregation on the `cell` column. The result is a `DataFrame`, and the row labels on the result are the names of the aggregation function.

```

table = indicators.agg({'life':['min', 'median'], 'cell':'median'})
print(table)

```

	life	cell
min	68.56	NaN
median	76.25	312.505

There is a NaN value here because we didn't request the minimum from the `cell` column

- If a data frame aggregation only has a single aggregation for any given column, the result can be a one-dimensional object (Series) with an entry for each of the columns being aggregated. In the following example, we compute the min aggregation for the `life` and `cell` columns.

```
avg = indicators.agg({'life': 'min', 'cell': 'min'})
print(avg)
```

```
life      68.56
cell      30.75
dtype: float64
```

Aggregating Selected Rows

- Often, we wish to characterize, through aggregation, a subset of a data frame. This “subset and aggregation” is also a critical piece in understanding the more general group partitioning and aggregation that we consider later.
- Consider the `topnames.csv` file (available on Teams). This dataset consists of the top male and female names, and the number (`count`) of babies with that name, from 1880 to 2018. The actual CSV file looks as follows:

```
year,sex,name,count
1880,Female,Mary,7065
1880,Male,John,9655
1881,Female,Mary,6919
1881,Male,John,8769
...
2017,Female,Emma,19800
2017,Male,Liam,18798
2018,Female,Emma,18688
2018,Male,Liam,19837
```

- Let’s read this data into a `DataFrame` and take a look at the first and last 10 rows (to compare which names were popular from different eras).

```
topnames = pd.read_csv("topnames.csv")
print(topnames.head(10))
print()
print(topnames.tail(10))
```

	year	sex	name	count
0	1880	Female	Mary	7065
1	1880	Male	John	9655
2	1881	Female	Mary	6919
3	1881	Male	John	8769
4	1882	Female	Mary	8148
5	1882	Male	John	9557
6	1883	Female	Mary	8012
7	1883	Male	John	8894
8	1884	Female	Mary	9217

John was the most popular male name in 1880 with 9655 babies given that name

9	1884	Male	John	9388
	year	sex	name	count
268	2014	Female	Emma	20936
269	2014	Male	Noah	19305
270	2015	Female	Emma	20455
271	2015	Male	Noah	19635
272	2016	Female	Emma	19496
273	2016	Male	Noah	19117
274	2017	Female	Emma	19800
275	2017	Male	Liam	18798
276	2018	Female	Emma	18688
277	2018	Male	Liam	19837

- Suppose we are interested in finding the minimum, maximum, and median of the `count` column, but only for rows where Mary was the top name. In the code below, we first create a subtable `mary_rows` that only contains the rows from `topnames` that correspond to Mary. When then use the `agg()` function to find the `min`, `max`, and `median` of the `count` column.

```
mary_rows = topnames[topnames['name'] == 'Mary']
table = mary_rows.agg({'count': ['min', 'max', 'median']})
print(table)
```

	count
min	6919.0
max	73985.0
median	54423.0

General Partitioning and `groupby`

- In mathematics, a *partition* of a set is a grouping of its elements into subsets in such a way that every element is included in exactly one subset.
- In data analytics, we may want to form a partition of a `DataFrame` using a *groupby* operation (pronounced “group by”). Each of these partitions then could have an aggregation performed upon it, and so, from each partition, we obtain a single *row* of values. In general, for the aggregation, we need to yield a result that is a single row per partition.
- The Pandas method for the partitioning/`groupby` operation is called `groupby()`. The first parameter of `groupby()` is the name of the column (or columns) to use for the partitioning. The result is a `DataFrameGroupBy` object, which is different from a `DataFrame` because it, in essence, is a set of data frames. We demonstrate below.

```
groupby_sex = topnames.groupby('sex')
groupby_name = topnames.groupby('name')

print("The number of groups by sex is: ", len(groupby_sex))
print("The number of groups by name is: ", len(groupby_name))
type(groupby_name)
```

The number of groups by sex is:	2
The number of groups by name is:	18

| pandas.core.groupby.generic.DataFrameGroupBy

- Thus, we can see that the partitioning by `sex` yields 2 groups, while the partitioning by `name` yields 18 groups (which is the same as the unique number of names from the data set). Finally, note that the type of object returned by `groupby()` is a `DataFrameGroupBy`.
- Suppose for the `groupby_name` partition, we wanted to compute the median and sum for `count`. This could be accomplished as follows.

```
namegroup = groupby_name.agg({'count': ['median', 'sum']})  
print(namegroup)
```

	count	
name	median	sum
Ashley	40965.5	81931
David	85929.0	85929
Emily	25104.0	294508
Emma	19648.0	118188
Isabella	22609.5	45219
Jacob	25339.0	370779
James	86224.0	1056228
Jennifer	57117.0	859209
Jessica	47884.0	397962
John	9032.5	861403
Liam	19317.5	38635
Linda	85723.5	508407
Lisa	53355.0	420572
Mary	54423.0	3098428
Michael	68117.5	3084824
Noah	19211.0	76314
Robert	60699.0	1041984
Sophia	21842.0	65378

Sorting

- We can change the order of the rows by sorting using the `sort_values()` method. To utilize this method, simply pass in the column name that you want to sort by.
- To see this in action, we will once again consider our `dogs` `DataFrame`.

```
sorted_dogs = dogs.sort_values("weight_kg")  
print(sorted_dogs)
```

	name	breed	color	height_cm	weight_kg	date_of_birth
5	Stella	Chihuahua	Tan	18	2	2015-04-20
3	Cooper	Schnauzer	Gray	49	17	2011-12-11
0	Bella	Labrador	Brown	56	24	2013-07-01
1	Charlie	Poodle	Black	43	24	2016-09-16
2	Lucy	ChowChow	Brown	46	24	2014-08-25
4	Max	Labrador	Black	59	29	2017-01-28
6	Bernie	St.Bernard	White	77	74	2018-02-27

- Note that the DataFrame is now sorted by `weight_kg` in *ascending* order. Further note that the `sort_value(.)` returns a sorted DataFrame but does not actually modify the DataFrame from which it was called.
- We can set the `ascending` argument to `False` so that it will sort the data in descending order.

```
sorted_dogs = dogs.sort_values("weight_kg", ascending=False)
print(sorted_dogs)
```

	name	breed	color	height_cm	weight_kg	date_of_birth
6	Bernie	St.Bernard	White	77	74	2018-02-27
4	Max	Labrador	Black	59	29	2017-01-28
0	Bella	Labrador	Brown	56	24	2013-07-01
1	Charlie	Poodle	Black	43	24	2016-09-16
2	Lucy	ChowChow	Brown	46	24	2014-08-25
3	Coooper	Schnauzer	Gray	49	17	2011-12-11
5	Stella	Chihuahua	Tan	18	2	2015-04-20

- We can sort by multiple variables by passing a list of columns to `sort_values(.)`. For example, we first sort by weight, then by height below.

```
sorted_dogs = dogs.sort_values(["weight_kg", "height_cm"])
print(sorted_dogs)
```

	name	breed	color	height_cm	weight_kg	date_of_birth
5	Stella	Chihuahua	Tan	18	2	2015-04-20
3	Coooper	Schnauzer	Gray	49	17	2011-12-11
1	Charlie	Poodle	Black	43	24	2016-09-16
2	Lucy	ChowChow	Brown	46	24	2014-08-25
0	Bella	Labrador	Brown	56	24	2013-07-01
4	Max	Labrador	Black	59	29	2017-01-28
6	Bernie	St.Bernard	White	77	74	2018-02-27

- Note that now Charlie, Lucy, and Bella are ordered from shortest to tallest, even though they weigh the same. To change the direction the values are sorted in, pass a list to the `ascending` argument indicating which direction sorting should be done for each variable:

```
sorted_dogs = dogs.sort_values(["weight_kg", "height_cm"],
                               ascending=[True, False])
print(sorted_dogs)
```

	name	breed	color	height_cm	weight_kg	date_of_birth
5	Stella	Chihuahua	Tan	18	2	2015-04-20
3	Coooper	Schnauzer	Gray	49	17	2011-12-11
0	Bella	Labrador	Brown	56	24	2013-07-01
2	Lucy	ChowChow	Brown	46	24	2014-08-25
1	Charlie	Poodle	Black	43	24	2016-09-16
4	Max	Labrador	Black	59	29	2017-01-28
6	Bernie	St.Bernard	White	77	74	2018-02-27

- Now Charlie, Lucy, and Bella are ordered from tallest to shortest.

Operations to Delete Columns and Rows

- The easiest way to delete a single column from a DataFrame is to use Python's `del` statement, using an argument specifying the single column.
- Before we start deleting columns, let's make a copy of our indicators DataFrame using the `copy()` method as follows.

```
ind2 = indicators.copy()
print(ind2)
```

	code	country	pop	gdp	life	cell
	CAN	Canada	36.26	1535.77	82.30	30.75
	CHN	China	1378.66	11199.15	76.25	1364.93
	IND	India	1324.17	2263.79	68.56	1127.81
	RUS	Russia	144.34	1283.16	71.59	229.13
	USA	United States	323.13	18624.47	78.69	395.88
	VNM	Vietnam	94.59	205.28	76.25	120.60

- We can remove the `cell` column from the `ind2` DataFrame using `del` as follows.

```
del ind2['cell']
print(ind2)
```

	code	country	pop	gdp	life
	CAN	Canada	36.26	1535.77	82.30
	CHN	China	1378.66	11199.15	76.25
	IND	India	1324.17	2263.79	68.56
	RUS	Russia	144.34	1283.16	71.59
	USA	United States	323.13	18624.47	78.69
	VNM	Vietnam	94.59	205.28	76.25

- Alternately, we can use the `pop()` method to delete a single column. This method deletes and modifies the data structure in place and returns the element that was deleted. We demonstrate below for the `gdp` column.

```
ind2 = indicators.copy()
gdp_series = ind2.pop('gdp')
print(ind2)
print(gdp_series)
```

	code	country	pop	life	cell
	CAN	Canada	36.26	82.30	30.75
	CHN	China	1378.66	76.25	1364.93
	IND	India	1324.17	68.56	1127.81
	RUS	Russia	144.34	71.59	229.13
	USA	United States	323.13	78.69	395.88
	VNM	Vietnam	94.59	76.25	120.60

Note that the gdp column has been dropped from ind2, but it has been stored in the Series gdp_series, which we can use later if needed.

	code	gdp
	CAN	1535.77
	CHN	11199.15
	IND	2263.79
	RUS	1283.16
	USA	18624.47
	VNM	205.28

Name: gdp, dtype: float64

Deleting Multiple Columns

- We can delete multiple columns using the `drop()` method of a `DataFrame`, where the first argument is a single column label, or a list of column labels. We demonstrate below where we delete columns `cell` and `life` from the original table.

```
ind2 = indicators.copy()
ind2.drop(['cell', 'life'], axis=1, inplace=True)
print(ind2)
```

	code	country	pop	gdp
	CAN	Canada	36.26	1535.77
	CHN	China	1378.66	11199.15
	IND	India	1324.17	2263.79
	RUS	Russia	144.34	1283.16
	USA	United States	323.13	18624.47
	VNM	Vietnam	94.59	205.28

- The `axis=1` argument indicates the dimension for what is to be dropped, in this case, 1 indicates the column (and an axis of 0 would indicate the row dimension). The `inplace=True` indicates that the method should *not* create a new `DataFrame`, but rather, drop the columns of the calling `DataFrame`. Note that `inplace=False` will create a new `DataFrame` and return the result but will leave the calling `DataFrame` alone.

Row Deletion

- We can use the `drop()` method to remove rows by specifying `axis=0`. We demonstrate below by dropping the rows corresponding to the USA and Russia.

```
ind2 = indicators.copy()
ind2.drop(['USA', 'RUS'], axis=0, inplace=True)
print(ind2)
```

	country	pop	gdp	life	cell
code					
CAN	Canada	36.26	1535.77	82.30	30.75
CHN	China	1378.66	11199.15	76.25	1364.93
IND	India	1324.17	2263.79	68.56	1127.81
VNM	Vietnam	94.59	205.28	76.25	120.60

Adding a Column

- One of the most common mutation operations is adding a new column to an existing DataFrame. We have already seen this earlier in the course, but we demonstrate it again below where we create a new column representing a 15% growth in GDP.

```
ind2 = indicators.copy()
ind2['growth'] = ind2.gdp*1.15
print(ind2)
```

	country	pop	gdp	life	cell	growth
code						
CAN	Canada	36.26	1535.77	82.30	30.75	1766.1355
CHN	China	1378.66	11199.15	76.25	1364.93	12879.0225
IND	India	1324.17	2263.79	68.56	1127.81	2603.3585
RUS	Russia	144.34	1283.16	71.59	229.13	1475.6340
USA	United States	323.13	18624.47	78.69	395.88	21418.1405
VNM	Vietnam	94.59	205.28	76.25	120.60	236.0720

- Since the `growth` column doesn't exist, Pandas creates the column and uses vectorization to fill in the data by multiplying the corresponding entries of the `gdp` column by 1.15.
- Another commonly used technique to create a new column is to use the `apply()` function to perform an operation on the elements of a vector.
- Recall that a lambda function is a small anonymous function that is defined without a name. For example,

```
to_caps = lambda s: s.upper()
```

will define a lambda function called `to_caps()` that takes a string and converts it to all capitals using the `upper()` string method.

- In the example below, we start with the original `indicators` DataFrame, drop the `life`, `cell`, `pop`, and `gdp` columns (for readability) and then use the lambda function `to_caps()` to create a new column with the country names in all capitals.

```
to_caps = lambda s: s.upper()

ind2 = indicators.copy().drop(['life', 'cell', 'pop', 'gdp'],
                             axis=1, inplace=False)
ind2['countryCaps'] = ind2['country'].apply(to_caps)
print(ind2)
```

	country	countryCaps
code		
CAN	Canada	CANADA
CHN	China	CHINA
IND	India	INDIA
RUS	Russia	RUSSIA
USA	United States	UNITED STATES
VNM	Vietnam	VIETNAM

Updating Columns

- Instead of creating an entirely new column in a data frame, we often want to change values in an existing column. To accomplish, we specify an existing column on the left-hand side of an assignment operator and a valid column-vector operation on the right-hand side. Below we demonstrate by updating the `life` column of the original `indicators` DataFrame.

```
ind2 = indicators.copy()
print('The original `life` column is')
print(ind2.life)
ind2.life = ind2.life + 0.5
print('\nThe updated `life` column is')
print(ind2.life)
```

```
The original `life` column is
code
CAN      82.30
CHN      76.25
IND      68.56
RUS      71.59
USA      78.69
VNM      76.25
Name: life, dtype: float64
```

```
The updated `life` column is
code
CAN      82.80
CHN      76.75
IND      69.06
RUS      72.09
USA      79.19
VNM      76.75
Name: life, dtype: float64
```

Combining Tables: Concatenating Along the Row Dimension

- Given two or more data frames, we may be interested in combining them into a single data frame that has some union or intersection of the rows, columns, and values of the source tables.
- Our examples below make use of the assumption that for the tables being combined, the source data frames represent different information. When combining along the row dimension, that implies that there is no intersection of value combination of the independent variable.
- In *normalized* tidy data, a **row index is meaningful** if the index is composed of the independent variable(s) of the data set. So, if a data set has one independent variable, the value of that variable is unique per row, and the remaining columns give the values of the dependent variables.
- Consider a subset of the `indicators` DataFrame consisting of the columns `country`, `pop`, `gdp`, and `life` for the rows `CHN`, `IND`, and `USA`:

```
ind1 = indicators.loc[['CHN', 'IND', 'USA'], : 'life']
print(ind1)
```

	country	pop	gdp	life
code				
CHN	China	1378.66	11199.15	76.25
IND	India	1324.17	2263.79	68.56
USA	United States	323.13	18624.47	78.69

- Now consider the new data frame below which contains values for the same columns as `ind1`, but for the countries `DEU` and `GBR`:

```
data2 = [[ 'Germany', 82.66, 3693.20, 80.99 ],
          ['United Kingdom', 66.06, 2637.87, 81.16]]

ind2 = pd.DataFrame(data2, columns=['country', 'pop', 'gdp', 'life'],
                    index=['DEU', 'GBR'])
print(ind2)
```

	country	pop	gdp	life
DEU	Germany	82.66	3693.20	80.99
GBR	United Kingdom	66.06	2637.87	81.16

- Thus, the **two DataFrames have the same columns, but different rows**. We can combine them using the `pd.concat()` function, where the first argument is a list of the data frames to be combined, followed by the `axis` to concatenate along (in our case, we are combining in the row dimension, so `axis=0`).


```
combined = pd.concat([ind1, ind2], axis=0)
print(combined)
```

	country	pop	gdp	life
CHN	China	1378.66	11199.15	76.25
IND	India	1324.17	2263.79	68.56
USA	United States	323.13	18624.47	78.69
DEU	Germany	82.66	3693.20	80.99
GBR	United Kingdom	66.06	2637.87	81.16

- Note that the order of the resulting data frame is based on the order of the data frames in the list.
- Consider the indicator data from 2015 for CHN, IND, and USA:

```
data2015 = [[ 'China', 1371.22, 11015.54, 76.09],
             ['India', 1310.15, 2103.59, 68.30],
             ['United States', 320.74, 18219.30, 78.69]]

indicators2015 = pd.DataFrame(data2015,
                              columns=['country', 'pop', 'gdp', 'life'],
                              index=['CHN', 'IND', 'USA'])

print(indicators2015)
```

	country	pop	gdp	life
CHN	China	1371.22	11015.54	76.09
IND	India	1310.15	2103.59	68.30
USA	United States	320.74	18219.30	78.69

- Now consider the indicator data from 2017 for CHN, IND, and USA:

```
data2017 = [[ 'China', 1386.40, 12143.49, 76.41],
             ['India', 1338.66, 2652.55, 68.80],
             ['United States', 325.15, 19485.39, 78.54]]

indicators2017 = pd.DataFrame(data2017,
                              columns=['country', 'pop', 'gdp', 'life'],
                              index=['CHN', 'IND', 'USA'])

print(indicators2017)
```

	country	pop	gdp	life
CHN	China	1386.40	12143.49	76.41
IND	India	1338.66	2652.55	68.80
USA	United States	325.15	19485.39	78.54

- If we combine in the same way as the last example, the result is a valid data frame, but one where the row labels have *duplicates*. While such non-uniqueness is allowed by Pandas, we no longer have a tidy data set. We demonstrate below.

```
combined = pd.concat([indicators2015, indicators2017], axis=0)
print(combined)
```

	country	pop	gdp	life
CHN	China	1371.22	11015.54	76.09
IND	India	1310.15	2103.59	68.30
USA	United States	320.74	18219.30	78.69
CHN	China	1386.40	12143.49	76.41
IND	India	1338.66	2652.55	68.80
USA	United States	325.15	19485.39	78.54

There are now duplicate row labels, and thus, this data is not tidy!

- The `concat()` function supports a `keys=` named parameter that allows us to specify a new, additional level for the row labels. The value is a list where each element specifies an index level value, so `keys=[2015, 2017]` will use an outer index of 2015 for the first data frame's rows and 2017 for the second data frame's rows.

```
combined = pd.concat([indicators2015, indicators2017], axis=0,
                      keys=[2015, 2017])
print(combined)
```

	country	pop	gdp	life
2015 CHN	China	1371.22	11015.54	76.09
IND	India	1310.15	2103.59	68.30
USA	United States	320.74	18219.30	78.69
2017 CHN	China	1386.40	12143.49	76.41
IND	India	1338.66	2652.55	68.80
USA	United States	325.15	19485.39	78.54

- This gives us the tidy two-level index with `year` and `code` as the levels. If we want the outer level to have a symbolic name, we can construct an `Index` object and specify both the list of outer level values as well as the name of the outer level index.

```
combined = pd.concat([indicators2015, indicators2017], axis=0,
                      keys=pd.Index([2015, 2017], name='year'))
print(combined)
```

year	country	pop	gdp	life
2015 CHN	China	1371.22	11015.54	76.09
IND	India	1310.15	2103.59	68.30
USA	United States	320.74	18219.30	78.69
2017 CHN	China	1386.40	12143.49	76.41
IND	India	1338.66	2652.55	68.80
USA	United States	325.15	19485.39	78.54

- Suppose we wanted to access all the data corresponding to 2017 from the combined `DataFrame`. We can accomplish this using `.loc[]` and Pandas slicing notation:

```
print(combined.loc[2017,:])
```

	country	pop	gdp	life
CHN	China	1386.40	12143.49	76.41
IND	India	1338.66	2652.55	68.80
USA	United States	325.15	19485.39	78.54

- To access a single value, say pop from USA in 2015, we can use a tuple within `.loc[]`:

```
print(combined.loc[(2015, 'USA'), 'pop'])
```

```
| 320.74
```

- Below we construct two data frames from the `topname` data: one for the years 2010 and 2011, and another from 2017 and 2018.

```
topnames1 = topnames[topnames['year'].isin([2010, 2011])]
print(topnames1)
print()
topnames2 = topnames[topnames['year'].isin([2017, 2018])]
print(topnames2)
```

	year	sex	name	count
260	2010	Female	Isabella	22913
261	2010	Male	Jacob	22127
262	2011	Female	Sophia	21842
263	2011	Male	Jacob	20371

	year	sex	name	count
274	2017	Female	Emma	19800
275	2017	Male	Liam	18798
276	2018	Female	Emma	18688
277	2018	Male	Liam	19837

- If we perform `concat()` we obtain the following:

```
table = pd.concat([topnames1, topnames2], axis=0)
print(table)
```

	year	sex	name	count
260	2010	Female	Isabella	22913
261	2010	Male	Jacob	22127
262	2011	Female	Sophia	21842
263	2011	Male	Jacob	20371
274	2017	Female	Emma	19800
275	2017	Male	Liam	18798
276	2018	Female	Emma	18688
277	2018	Male	Liam	19837

- Note that the row indices run from 260-263, and then from 274-277. This could potentially cause errors, especially if we are performing a `.loc[]` operation. We can reindex our data frame using the `reset_index()` method:

```
table.reset_index(inplace=True)
print(table)
```

By setting
`inplace=True` we tell
Python to update the
calling DataFrame.
Extremely useful!

	index	year	sex	name	count
0	260	2010	Female	Isabella	22913
1	261	2010	Male	Jacob	22127
2	262	2011	Female	Sophia	21842
3	263	2011	Male	Jacob	20371
4	274	2017	Female	Emma	19800
5	275	2017	Male	Liam	18798
6	276	2018	Female	Emma	18688
7	277	2018	Male	Liam	19837

Combining Tables: Concatenating Along the Column Dimension

- Here, we assume that the two source data frames have the *same rows* and an entirely *different set of columns*. In tidy data, the stipulation of the same rows means that the values of the independent variables are the same, and here we assume that they are incorporated into a meaningful index, and the stipulation of different columns means the data frames have different dependent variables
- Below we define a subtable called `cols1` that includes the `country` and `pop` columns for IND, CHN, and USA from the `indicators` DataFrame. In addition, we define a DataFrame `cols2` that includes the columns `imports` and `exports` for the same countries (rows).

```
cols1 = indicators.loc[['IND', 'CHN', 'USA'], ['country', 'pop']]

imports_exportsDoL = {
    'imports': [2241.66, 392.23, 1601.76],
    'exports': [1504.57, 266.16, 2280.54]}

codes = pd.Index(['USA', 'IND', 'CHN'],
                  name='code')

cols2 = pd.DataFrame(imports_exportsDoL, index=codes)

print(cols1)
print()
print(cols2)
```

	country	pop
code		
IND	India	1324.17
CHN	China	1378.66
USA	United States	323.13

	imports	exports
code		
USA	2241.66	1504.57
IND	392.23	266.16
CHN	1601.76	2280.54

- Note that both `cols1` and `cols2` have an index defined by `code`. We can combine them along the column dimension by the argument `axis=1` as follows.

```
table = pd.concat([cols1, cols2], axis=1)
print(table)
```

	country	pop	imports	exports
code				
IND	India	1324.17	392.23	266.16
CHN	China	1378.66	1601.76	2280.54
USA	United States	323.13	2241.66	1504.57

- Note that the order of the columns was dictated by the order of the first DataFrame.

Joining Data Frames

- Pandas provides more powerful tools when combining data frames in more complicated ways than just along the row or column dimensions.
- One variation, `join()`, is a DataFrame method and uses the row label index for matching rows between the source frames. The other variation, `merge()`, is a function of the Pandas package and can use any column (or index level) for matching rows between the source frames.
- The combined frame has columns from both the original source frames with values populated based on the matching rows. If the two frames have columns with the same column name, the join/merge will include both, with the column names modified to distinguish the source frame of the overlapping column.
- Consider the two data frames below.

```
imports_exportsDoL = {
    'imports': [457.46, 643.52, 2342.67],
    'exports': [418.86, 441.11, 1545.61]}

codes = pd.Index(['CAN', 'GBR', 'USA'], name='code')

join1 = pd.DataFrame(imports_exportsDoL, index=codes)

country_landDoL = {
    'country': ['Belgium', 'United Kingdom', 'United States', 'Vietnam'],
    'land': [30280.0, 241930.0, 9147420.0, 310070.0]}

codes = pd.Index(['BEL', 'GBR', 'USA', 'VNM'], name='code')

join2 = pd.DataFrame(country_landDoL, index=codes)

print(join1)
print()
print(join2)
```

	imports	exports	
code			
CAN	457.46	418.86	← join1 data frame
GBR	643.52	441.11	
USA	2342.67	1545.61	

	country	land	
code			
BEL	Belgium	30280.0	← join2 data frame
GBR	United Kingdom	241930.0	
USA	United States	9147420.0	
VNM	Vietnam	310070.0	

- Suppose we want all rows of `join1` to be represented in the combined result that includes `country` and `land`, but if a match is not found in `join2`, we just fill with missing values for `country` and `land`. If `join1` is the first (or left-hand) frame written as we combine, this type of combination is called a **left join**.
- We invoke the `join()` method on the first data frame and specify a second argument of the “right” data frame. The named parameter `how=` is passed a string “left” to specify a left join, as illustrated below.

```
table = join1.join(join2, how="left")
print(table)
```

	imports	exports	country	land
code				
CAN	457.46	418.86	NaN	NaN
GBR	643.52	441.11	United Kingdom	241930.0
USA	2342.67	1545.61	United States	9147420.0

- Note the appropriate values for the matched rows, and missing values where the right-hand frame did not have a match to correspond to the left frame. In this situation, we think of the left frame as being dominant, and taking whatever values it can from the right frame, to enrich rows present in the left frame.
- If the desired analysis requires that the combined result only has rows where a row label index matches from both frames, this is called an **inner join**. It is a form of intersection, where the intersection is defined by matching just the value of the common row label index and disregarding any/all other column values. This inner join simply changes the `how=` named parameter argument to “inner”:

```
table = join1.join(join2, how="inner")
print(table)
```

	imports	exports	country	land
code				
GBR	643.52	441.11	United Kingdom	241930.0
USA	2342.67	1545.61	United States	9147420.0

- The result includes just the two rows where `code` is in common between the two frames, namely GBR and USA. In this case, we think of the new table as sitting between the two old tables, and drawing from both sides, whenever the same code is present in both.
- If we were to join the two tables in a different order, making `join2` be the lefthand frame and `join1` be the right-hand frame, then, in a left join, the result would have all the rows from `join2` with the `join2` columns of `country` and `land` and would fill in with missing values for the columns in `join1` (`imports` and `exports`) where there was no corresponding row in `join1`. In this case, `join2` is the dominant table.

```
table = join2.join(join1, how="left")
print(table)
```

	code	country	land	imports	exports
	BEL	Belgium	30280.0	NaN	NaN
	GBR	United Kingdom	241930.0	643.52	441.11
	USA	United States	9147420.0	2342.67	1545.61
	VNM	Vietnam	310070.0	NaN	NaN

- Rows BEL and VNM are present in `join2` but not in `join1`, so these are the ones that are augmented with missing values for the `join1` columns.

Merging Data Frames

- Sometimes, we wish to combine two tables based on common values between two tables, but the values are in a regular column and not part of an index.
- Recall that the `topnames` data frame has, by year and by sex, the name and count of the top baby names. The original `indicators` data frame has, among other things, the population for each of the countries, including the population for the USA.
- Suppose we wanted to create a new data frame based primarily on the `topnames` data set, but where we augment each row with the US population for that year. That could be used in an analysis, to divide the count of applicants by the population for that year to be able to see what percentage of the population is represented. This could give a fairer comparison as we try to compare between years.
- Below, we load in the original `indicators` data frame, which contains economic indicator data for 207 countries from 1960-2018. Then we create a new data frame called `us_pop`, that contains the population (in millions) for the USA from 1960 to 2017 by extracting the data from `indicators`. Can you follow how the data is extracted?

```

table = indicators[(indicators['code'] == 'USA') &
                    (indicators['year'] >= 1960) &
                    (indicators['year'] <= 2017)].reset_index()

us_pop = table[['year', 'pop']]
print(us_pop.head())

```

	year	pop
0	1960	180.67
1	1961	183.69
2	1962	186.54
3	1963	189.24
4	1964	191.89

- Recall the topnames data frame:

```
print(topnames.head())
```

	year	sex	name	count
0	1880	Female	Mary	7065
1	1880	Male	John	9655
2	1881	Female	Mary	6919
3	1881	Male	John	8769
4	1882	Female	Mary	8148

- The `merge()` function has, as its first two arguments, the two DataFrame objects to be combined, where the first argument is considered the “left” and the second argument is considered the “right.” The `on=` named argument specifies the name of a column expected to exist in both data frames and to be used for matching values. The `how=` named parameter allows us to specify the logical equivalent of an inner join, a left join, or a right join.
- Below, we illustrate a left join/merge. Since `topnames` is the “left,” this will give us all the rows of `topnames` and will add the `pop` column from matching rows in `us_pop`. For those rows and years where `us_pop` does not have population data (i.e., those years before 1960), the new data frame `mergel` has NaN indicating missing data.

```

mergel = pd.merge(topnames, us_pop, on='year', how='left')
print(mergel.head())
print()
print(mergel[(mergel['year'] >= 1960) & (mergel['year'] <= 2017)].head())

```

	year	sex	name	count	pop
0	1880	Female	Mary	7065	NaN
1	1880	Male	John	9655	NaN
2	1881	Female	Mary	6919	NaN
3	1881	Male	John	8769	NaN
4	1882	Female	Mary	8148	NaN

mergel for the years 1880-1882—note the NaN values for pop.

	year	sex	name	count	pop
160	1960	Female	Mary	51475	180.67
161	1960	Male	David	85929	180.67
162	1961	Female	Mary	47680	183.69
163	1961	Male	Michael	86917	183.69
164	1962	Female	Lisa	46078	186.54

mergel for the years 1960-1962—note the pop contains the merged information.

- Now we illustrate an “inner” join/merge. Here, the result will only have rows where the year has common values from both data frames. So, this result will effectively prune the `topnames` portion to the years after 1960.

```
merge2 = pd.merge(topnames, us_pop, on='year', how='inner')
print(merge2)
```

	year	sex	name	count	pop
0	1960	Female	Mary	51475	180.67
1	1960	Male	David	85929	180.67
2	1961	Female	Mary	47680	183.69
3	1961	Male	Michael	86917	183.69
4	1962	Female	Lisa	46078	186.54
..
109	2014	Male	Noah	19305	318.56
110	2015	Female	Emma	20455	320.90
111	2015	Male	Noah	19635	320.90
112	2016	Female	Emma	19496	323.13
113	2016	Male	Noah	19117	323.13

[114 rows x 5 columns]

Missing Data Handling

- We end this section with a brief discussion of the issue with missing data. Missing data can occur when:
 - Individuals do not respond to all questions on a survey.
 - Countries fail to maintain or report all their data to the World Bank.
 - An organization keeping personnel records does not know where everyone lives.
 - Laws prevent healthcare providers from disclosing certain types of data.
 - Different users of social media select different privacy settings, resulting in some individuals having only some of their data publicly visible.
 - Many other situations analogous to these.
- In the real world, there is no true standard for how missing data is encoded. It may be coded as a string such as “N/A”, “Not Applicable”, “NA”, etc. Thankfully, information on how missing data is coded is typically contained in the *metadata* associated with the data set, sometimes called a *codebook*.
- Once we know how missing data is denoted, we should replace that coding with a special type. In Pandas, this type is denoted `nan` and is displayed as `NaN`. Please refer back to our section on Numpy to review how `NaN` values are treated.

Creating New Columns

- It is common to create new columns derived from existing columns. Creating and adding new columns can go by many names, including *mutating* a DataFrame, *transforming* a DataFrame, and *feature engineering*.
- Suppose we wanted to create a new column that has each dog's height in meters instead of centimeters. We can accomplish this with the command

```
dogs["height_m"] = dogs["height_cm"]/100.
```

On the left-hand side of the equals, we use square brackets with the name of the new column we want to create. On the right-hand side, we have the calculation. Note that we are utilizing vectorization! Below we show this in action.

```
dogs["height_m"] = dogs["height_cm"]/100
print(dogs)
```

	name	breed	color	height_cm	weight_kg	date_of_birth	height_m
0	Bella	Labrador	Brown	56	24	2013-07-01	0.56
1	Charlie	Poodle	Black	43	24	2016-09-16	0.43
2	Lucy	ChowChow	Brown	46	24	2014-08-25	0.46
3	Coooper	Schnauzer	Gray	49	17	2011-12-11	0.49
4	Max	Labrador	Black	59	29	2017-01-28	0.59
5	Stella	Chihuahua	Tan	18	2	2015-04-20	0.18
6	Bernie	St.Bernard	White	77	74	2018-02-27	0.77

- Notice that both the existing column and the new column we just created are in the DataFrame.
- Now let's add a column that has the body mass index (BMI) of the dogs. Recall that BMI is calculated using the formula

$$\text{BMI} = (\text{weight in kg}) / (\text{height in m})^2.$$

We can once again use vectorization to create the column as follows.

```
dogs['bmi'] = dogs['weight_kg'] / dogs['height_m']**2
print(dogs[['name', 'breed', 'color', 'date_of_birth', 'bmi']])
```

	name	breed	color	date_of_birth	bmi
0	Bella	Labrador	Brown	2013-07-01	76.530612
1	Charlie	Poodle	Black	2016-09-16	129.799892
2	Lucy	ChowChow	Brown	2014-08-25	113.421550
3	Coooper	Schnauzer	Gray	2011-12-11	70.803832
4	Max	Labrador	Black	2017-01-28	83.309394
5	Stella	Chihuahua	Tan	2015-04-20	61.728395
6	Bernie	St.Bernard	White	2018-02-27	124.810255

- Notice that in the code above we only choose to print some of the columns.